

CS 5/7320
Artificial Intelligence

Search with Uncertainty

AIMA Chapters 4.3-4.5

Slides by Michael Hahsler
with figures from the AIMA textbook



This work is licensed under a [Creative Commons Attribution-ShareAlike 4.0 International License](https://creativecommons.org/licenses/by-sa/4.0/).



Types of uncertainty we consider for now*



Nondeterministic Actions:

Outcome of an action in a state is uncertain.



No observations:

Sensorless problem



Partially observable environments:

The agent does not know in what state it is.



Unknown environments:

Online search

* we will quantify uncertainty with probabilities later.

Consequence of Uncertainty

- **Remember:** The solution for the known maze was a fixed **sequence of actions** from start to goal.
- **With uncertainty:** Solution is not a precomputed sequence, but a
conditional plan (a strategy or policy)
that depends on percepts.

A dynamic background image showing a bright yellow powder or smoke explosion against a solid black background. The explosion originates from the right side and spreads towards the left, with many fine particles visible in the air.

Nondeterministic Actions

Nondeterministic Actions

Outcome of actions in the environment is nondeterministic = **transition model need to describe uncertainty**

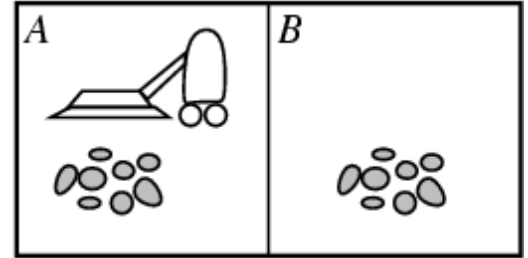
Example transition:

$$Results(s_1, a) = \{s_2, s_4, s_5\}$$

i.e., action a in s_1 can lead to one of several states.

The set of possible states $b = \{s_2, s_4, s_5\}$ is called a **belief state** of the agent.

Example: Erratic Vacuum World

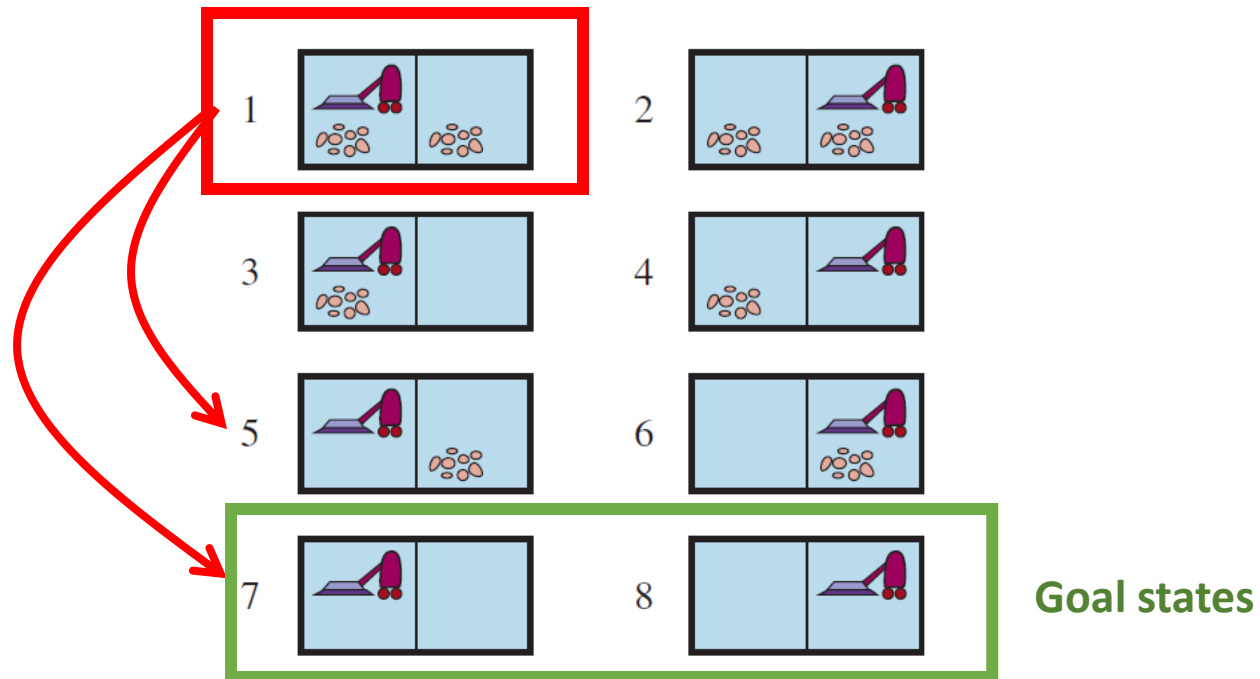


Regular vacuum world, but the action '**suck**' is more powerful and nondeterministic:

- a) **On a dirty square:** cleans the square and sometimes cleans dirt on adjacent squares as well.
- b) **On a clean square:** sometimes deposits some dirt on the square.

Example: Erratic Vacuum World

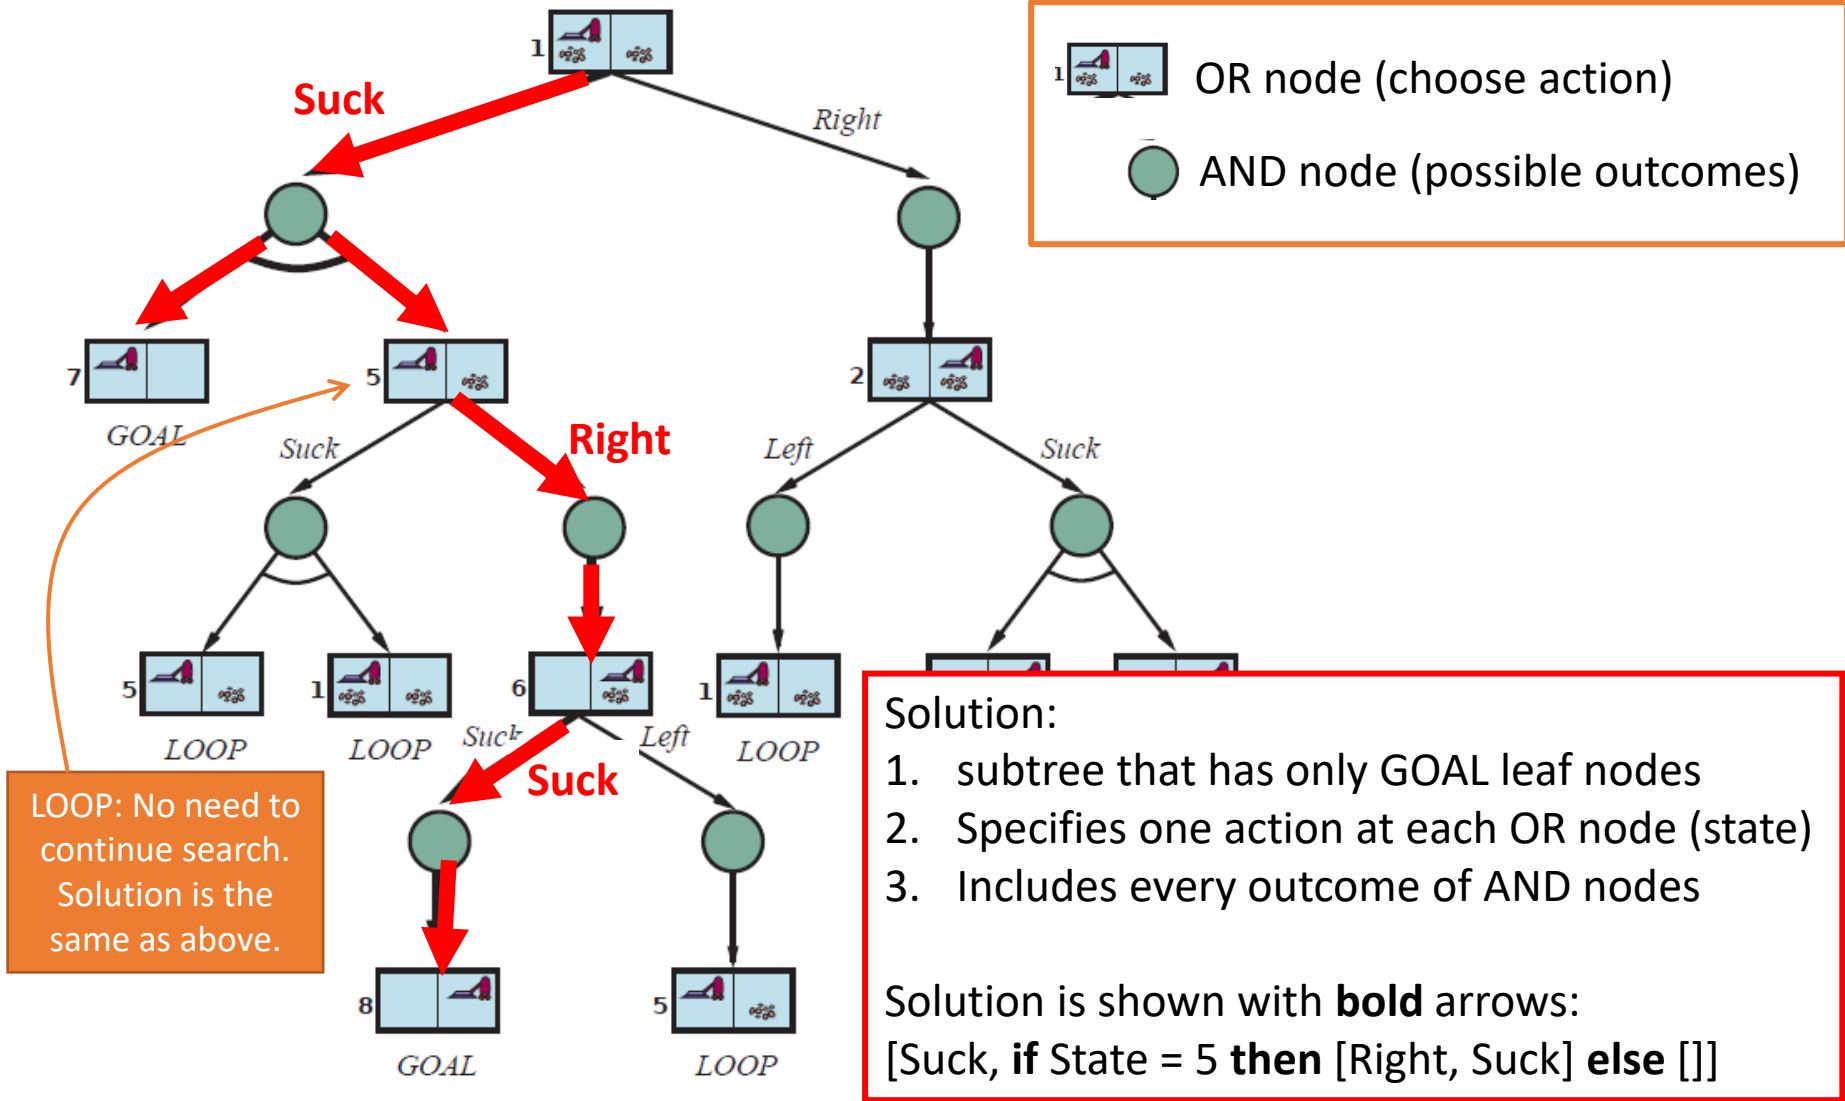
$Results(1, Suck) = \{5, 7\}$



We need a conditional plan

[Suck, **if** State = 5 **then** [Right, Suck] **else** []]

AND-OR Search Tree



AND-OR Recursive DFS Algorithm

= nested If-then-else statements

function AND-OR-SEARCH(*problem*) **returns** a conditional plan, or *failure*
 return OR-SEARCH(*problem*, *problem*.INITIAL, [])

path is only maintained for cycle checking!

function OR-SEARCH(*problem*, *state*, *path*) **returns** a conditional plan, or *failure*

if *problem*.IS-GOAL(*state*) **then return** the empty plan

if IS-CYCLE(*path*) **then return** *failure*

// don't follow loops using path.

for each *action* **in** *problem*.ACTIONS(*state*) **do**

// try all possible actions

plan \leftarrow AND-SEARCH(*problem*, RESULTS(*state*, *action*), [*state*] + *path*)

if *plan* \neq *failure* **then return** [*action*] + *plan*

return *failure*

function AND-SEARCH(*problem*, *states*, *path*) **returns** a conditional plan, or *failure*

for each s_i **in** *states* **do**

// try all possible outcomes, none can fail!

*plan*_{*i*} \leftarrow OR-SEARCH(*problem*, s_i , *path*)

// (= belief state)

if *plan*_{*i*} = *failure* **then return** *failure*

return [if s_1 **then** *plan*₁ **else if** s_2 **then** *plan*₂ **else** ... **if** s_{n-1} **then** *plan* _{$n-1$} **else** *plan* _{n}]

BFS and A* search can also be used to search an AND-OR tree.

Search with no
Observations



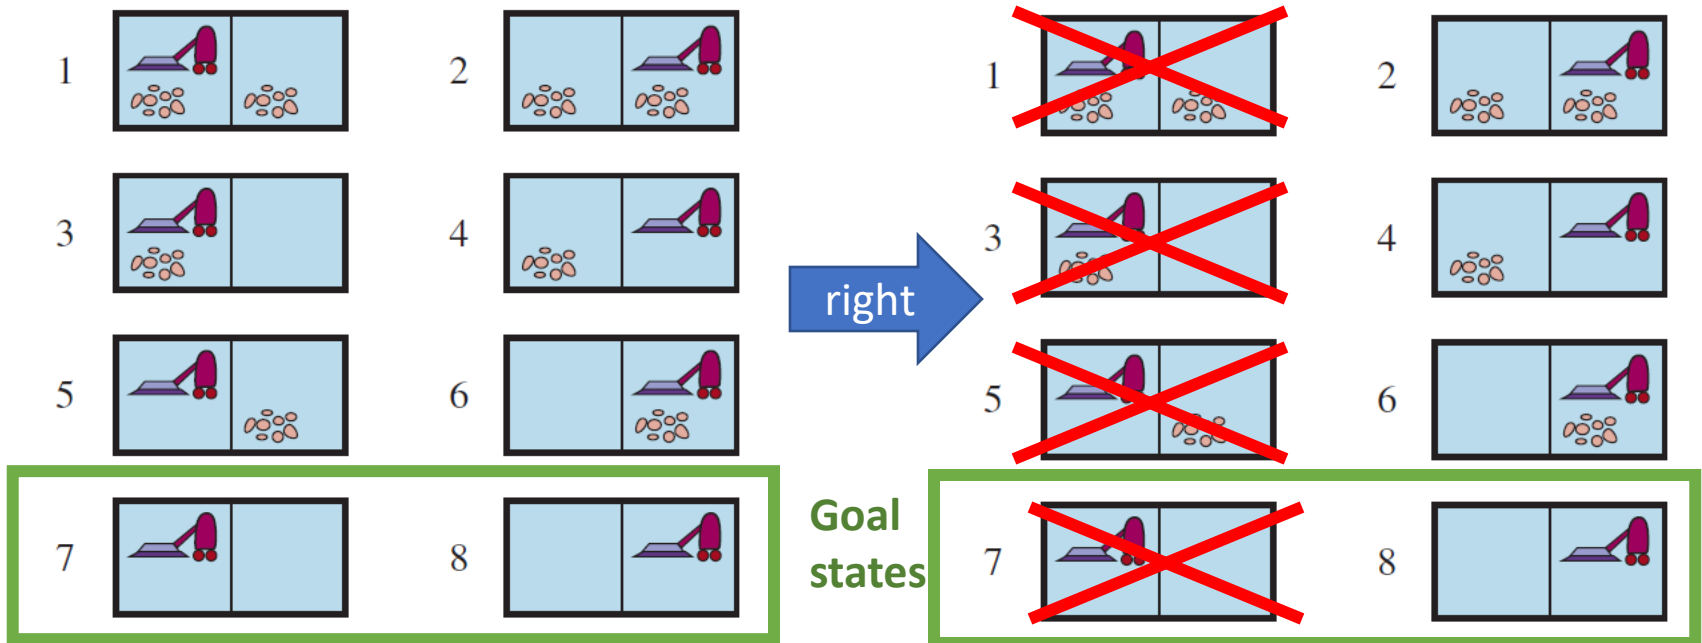
No Observations

- Sensorless problem = conformant problem
- **Example:** Doctor prescribes a broad-band antibiotic instead of performing time-consuming blood work for a more specific antibiotic. This saves time and money.
- **Basic idea:** Find a solution (a sequence of actions) that works from any state.

Actions to Coerce the World into States

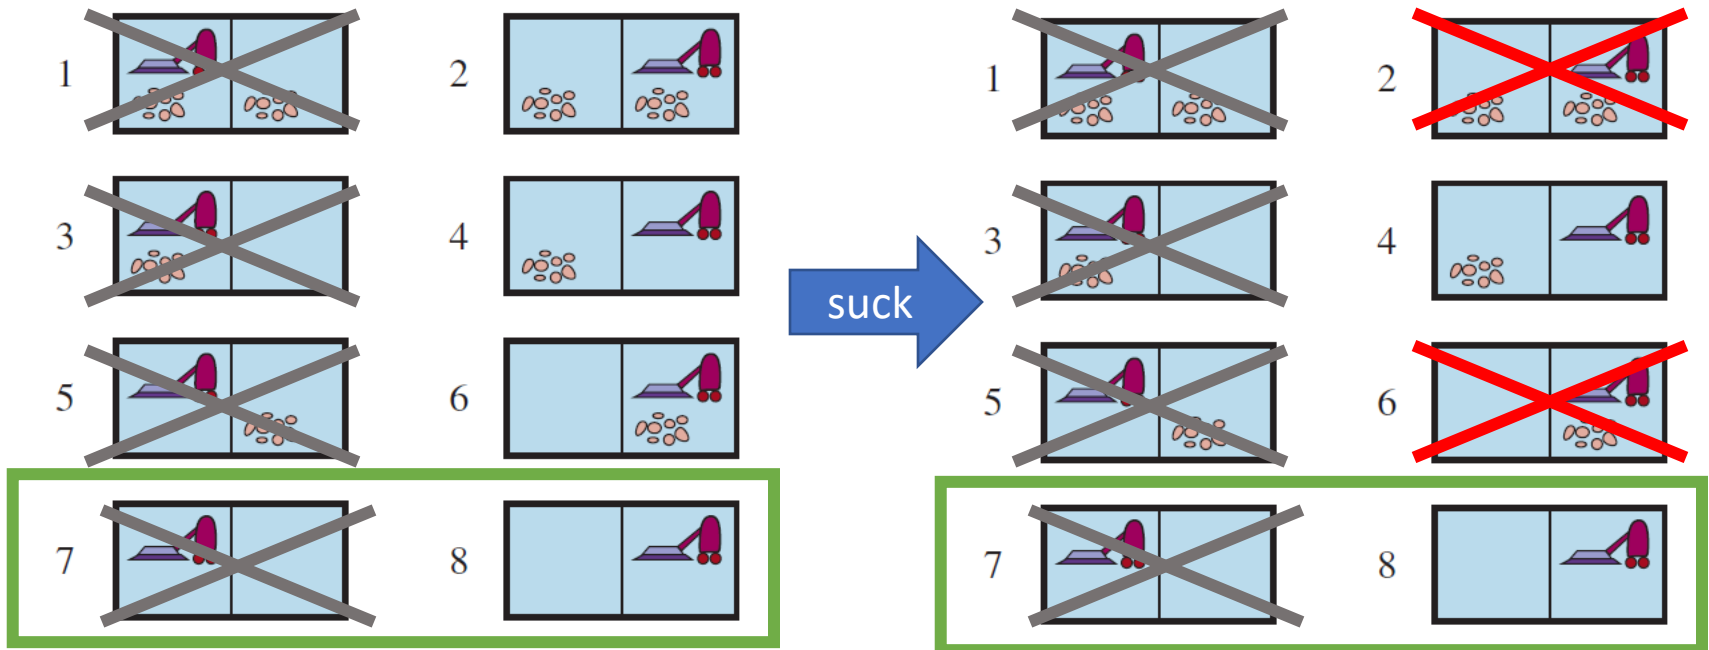
- Actions can reduce the number of possible states.
- **Example:** Deterministic vacuum world. Agent does not know its position and the dirt distribution.

Initial belief state {1,2,3,4,5,6,7,8}



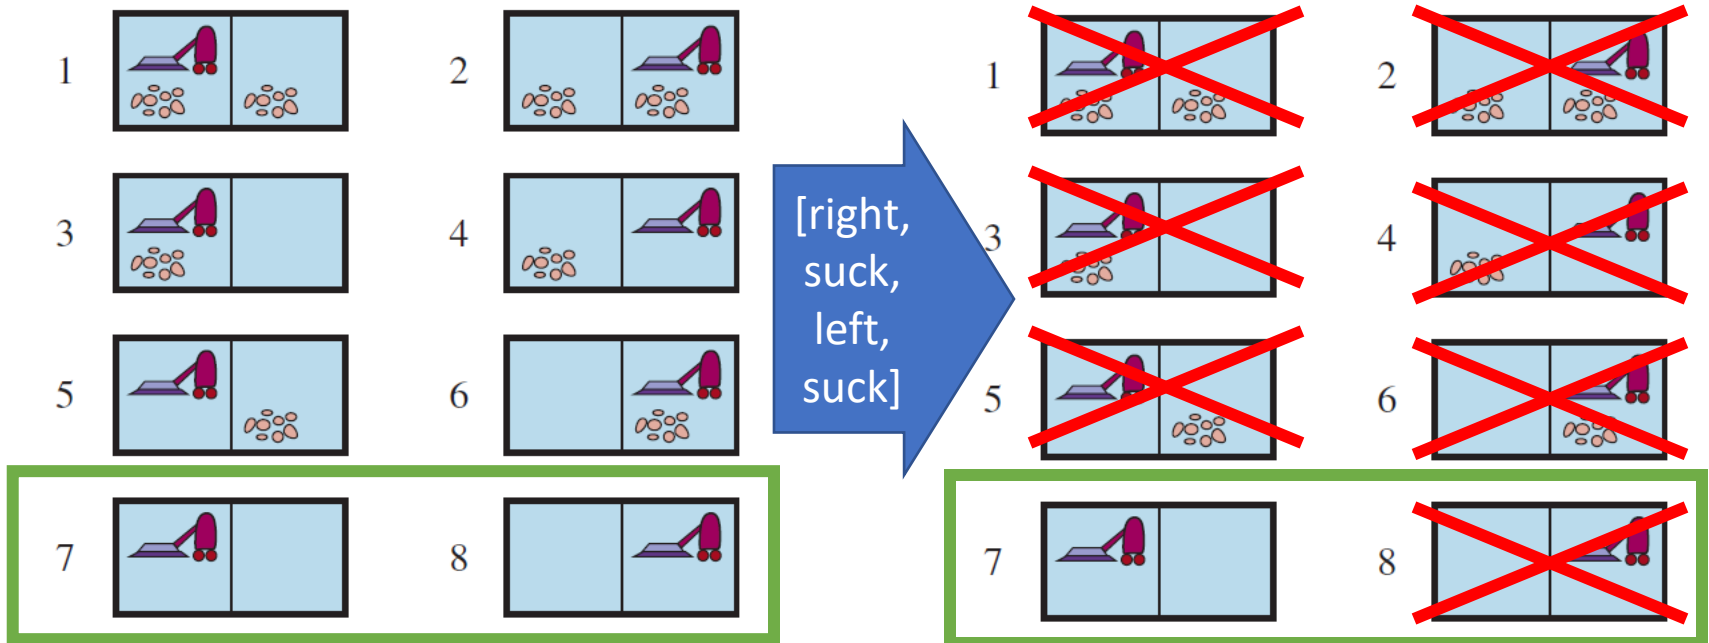
Actions to Coerce the World into States

- Actions can reduce the number of possible states.
- **Example:** Deterministic vacuum world. Agent does not know its position and the dirt distribution.



Actions to Coerce the World into States

- The action sequence [right, suck, left, suck] coerces the world into the goal state 7.
- This works from any initial state and there is no need for a conditional plan!



Find the Solution Sequence

Note: State space size makes this impractical for larger problems!

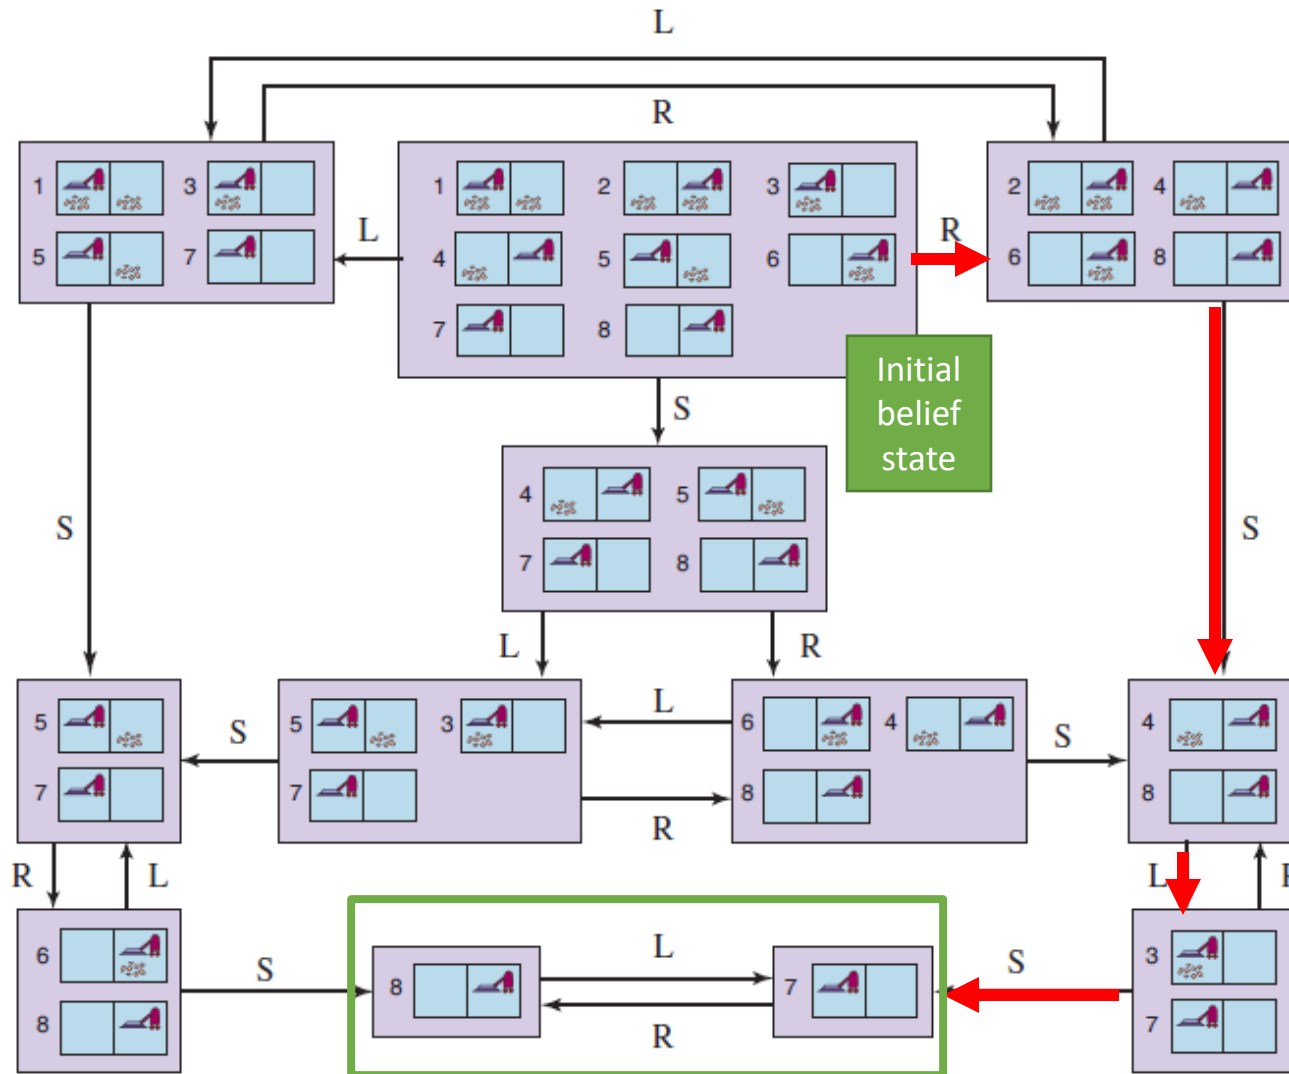
Use regular search (DFS, BFS, A*) for the problem:

- **States:** All belief states (=powerset \mathcal{P}_s of states of size 2^N for N states)
- **Initial state:** Often the belief state consisting of all states.
- **Actions:** Actions of a belief state are the union of the possible actions for all the states it contains.
- **Transition model:** $b' = Results(b, a) = \{s' : s' = Result(s, a) \text{ and } s \in b\}$
- **Goal test:** Are all states in the belief state goal states?
- **Simplifying property:** If a belief state (e.g., $b_1 = \{1,2,3,4,5\}$) is solvable (i.e., there is a sequence of actions that coerce all states to only goal states), then belief states that are subsets (e.g., $b_2 = \{2,5\}$) are also solved using the same action sequence. Used to prune the search tree.

Other approach:

- **Incremental belief-state search.** Generate a solution that works for one state and check if it also works for all other states.

Reachable belief-state space for the deterministic, sensorless vacuum world



Solution Sequence:
[right, suck, left, suck]

Size of the belief state space:

$$\mathcal{P}_s = 2^N = 2^8 = 256$$

Only as small fraction is reachable.

A close-up photograph of various colorful, 3D geometric shapes, including letters and numbers, scattered on a blue surface. The shapes are made of a matte plastic material and come in primary colors: red, yellow, and blue. Some shapes are in sharp focus, while others are blurred in the background, creating a sense of depth. The shapes include a large red 'V', a yellow '3', a red '2', a yellow '8', a red 'X', a yellow '0', a yellow '1', a blue '0', and a green 'C'. The text "Partially Observable Environments" is overlaid in the center in a white, sans-serif font.

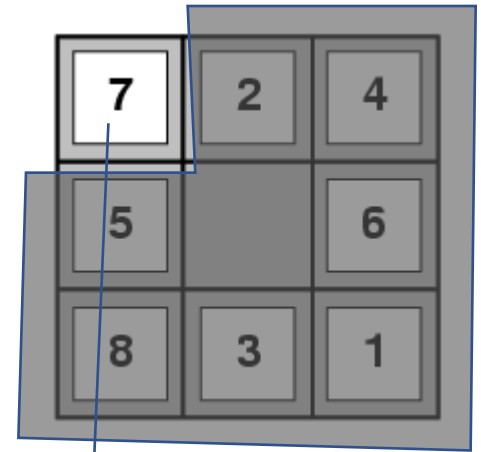
Partially Observable Environments

Percepts

- Many problems cannot be solved efficiently without sensing (e.g., 8-puzzle).
- We need to be able to at least see one square.

Percept function: $Percept(s)$

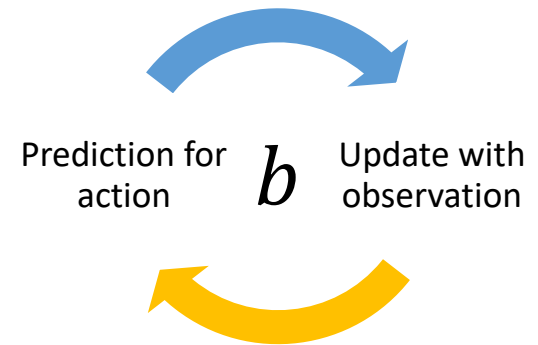
- **Fully observable:** $Percept(s) = s$
- **Sensorless:** $Percept(s) = null$
- **Partially observable:** $Percept(s) = o$



$Percept(s) = \text{Tile7}$

Problem: Many states can produce the same percept!

Prediction and Update



Assume we have a belief state b (i.e., the set of states we could be in)

Prediction: Compute new belief state that results from action a .

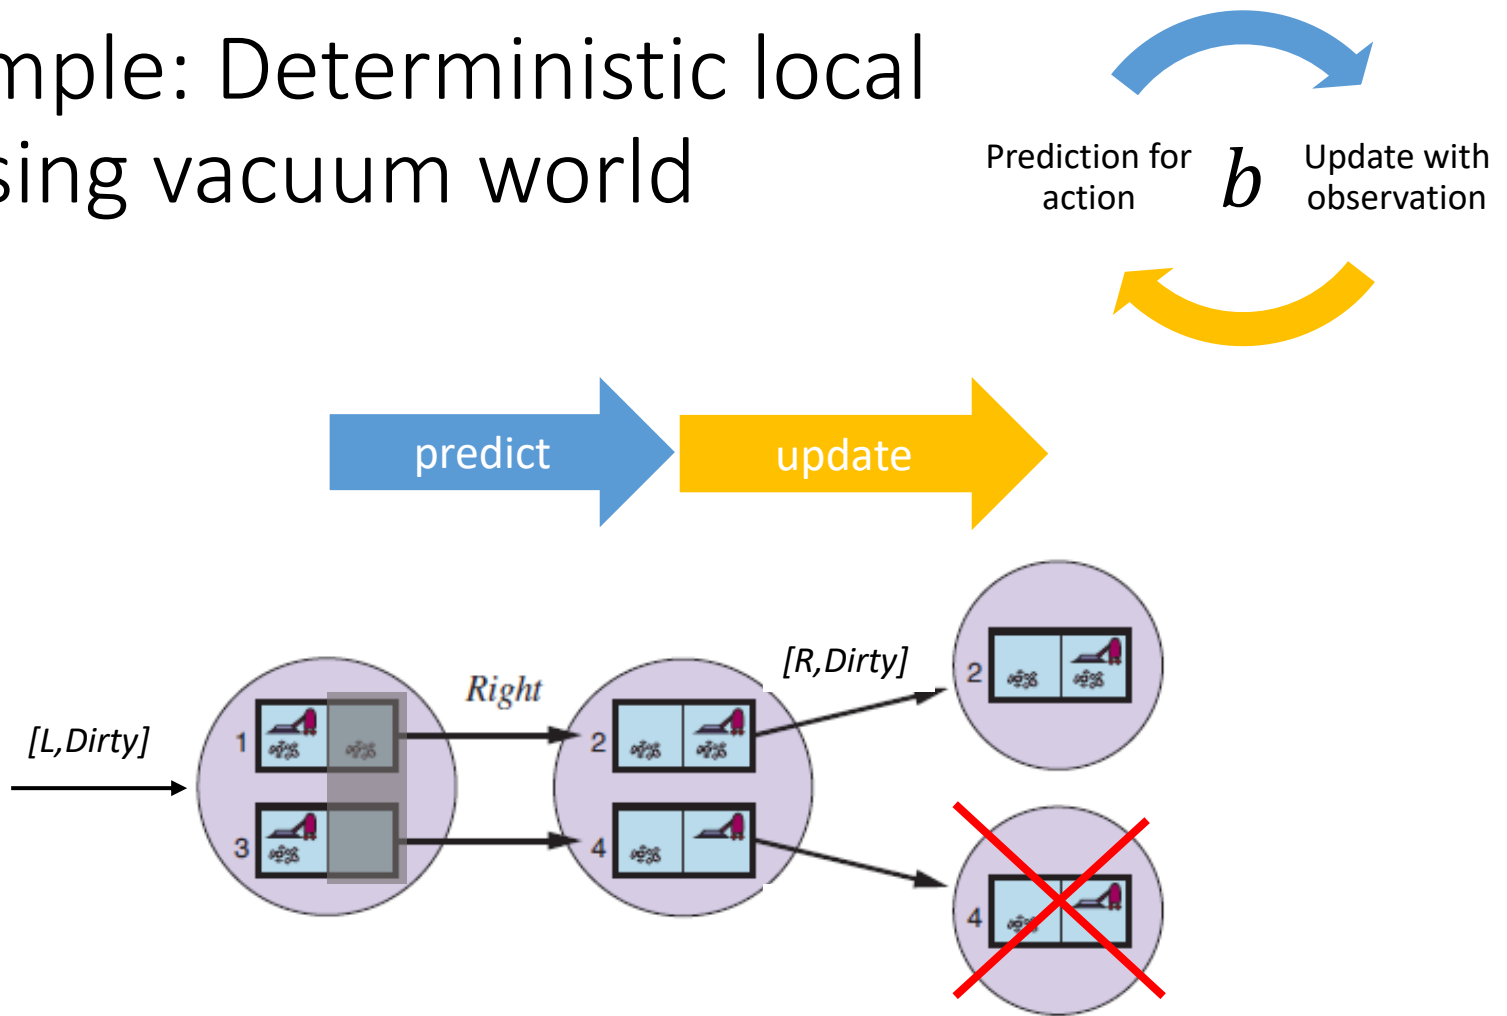
$$\hat{b} = \text{Predict}(b, a)$$

Update: Filter states that are consistent with new observation o .

$$b_o = \text{Update}(\hat{b}, o) = \{s : o = \text{Percept}(s) \text{ and } s \in \hat{b}\}$$

Both steps in one: $b \leftarrow \text{Update}(\text{Predict}(b, a), o)$

Example: Deterministic local sensing vacuum world

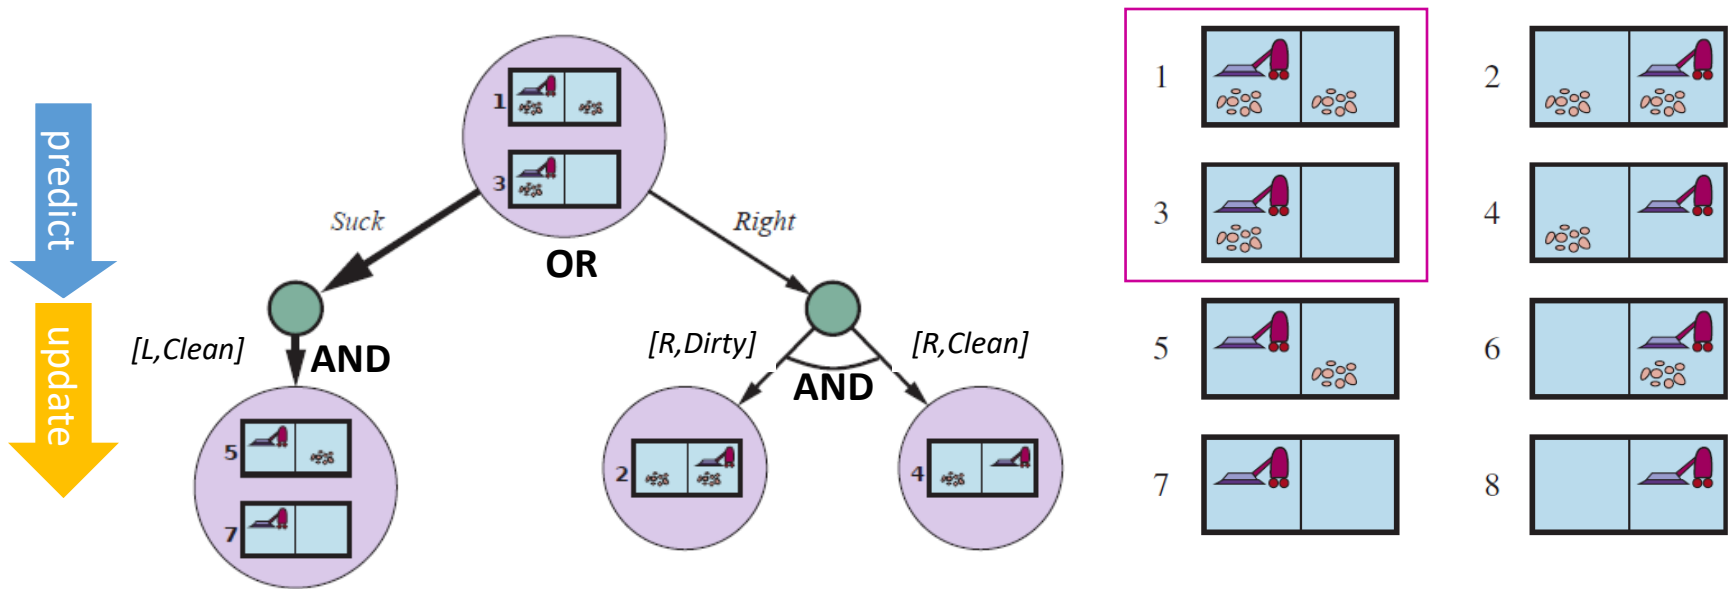


$$b \leftarrow \text{Update}(\text{Predict}(b, a), o)$$

$$\text{Update}(\text{Predict}(\{1, 3\}, \text{Right}), [R, \text{Dirty}]) = \{2\}$$

Solving Partially Observable Problems

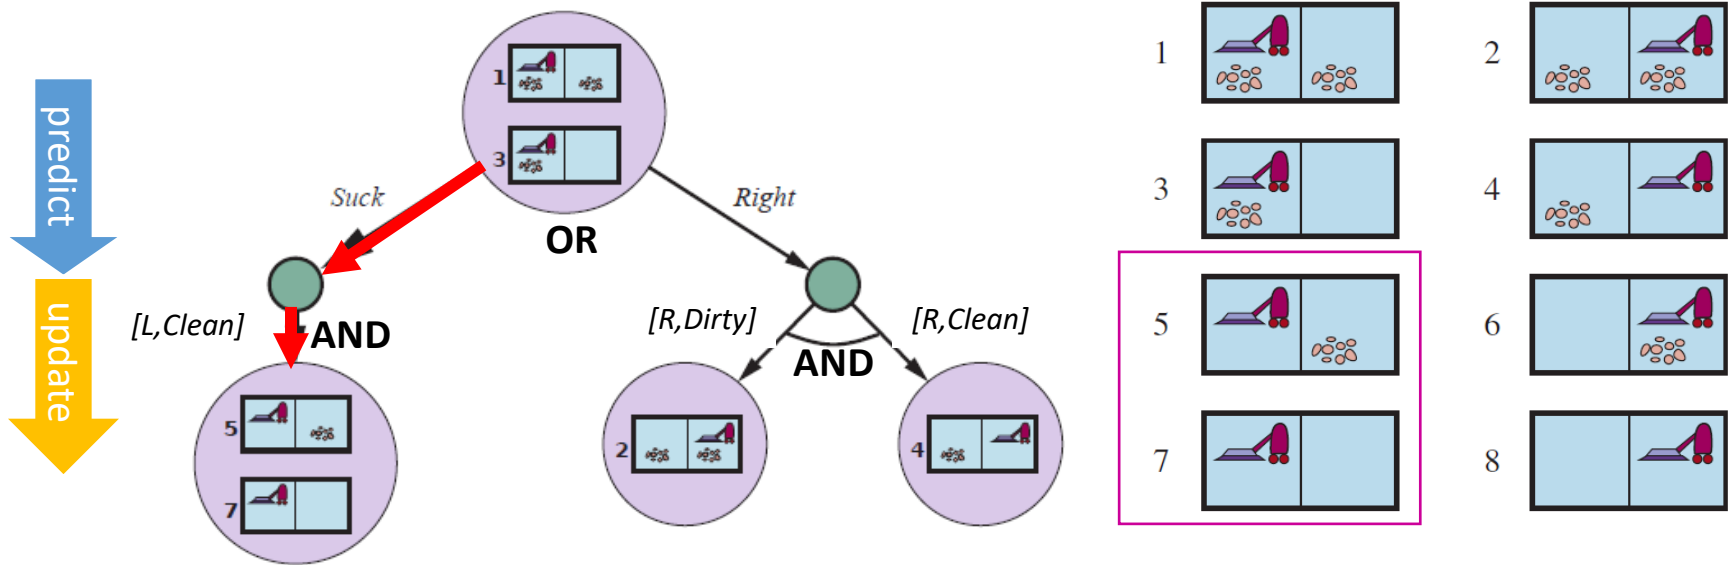
Use an AND-OR tree to create a conditional plan



Solution: $[Suck, Right, \text{if } b = \{6\} \text{ then } Suck \text{ else } []]$

Solving Partially Observable Problems

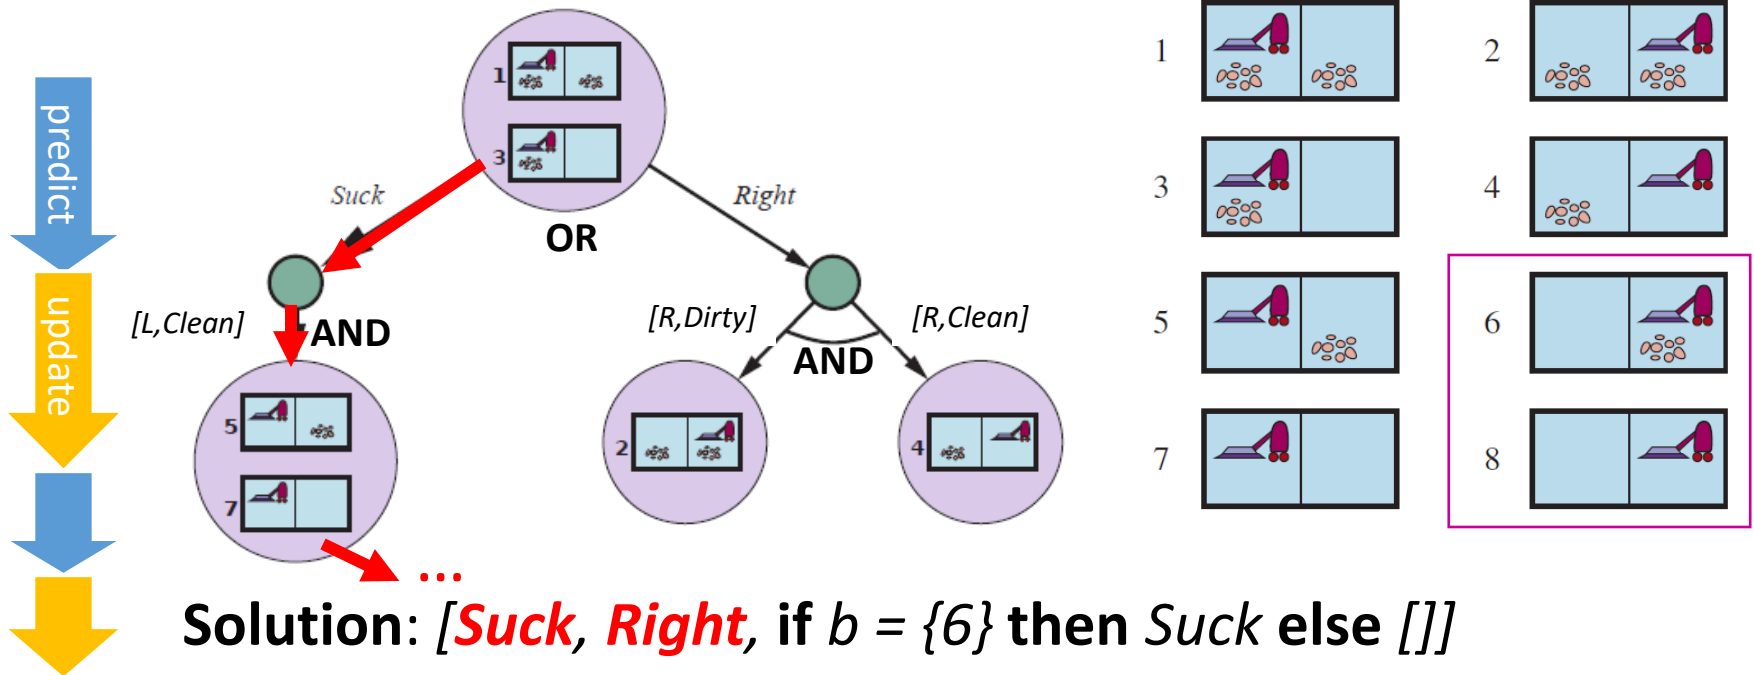
Use an AND-OR tree to create a conditional plan



Solution: [*Suck*, Right, if $b = \{6\}$ then *Suck* else []]

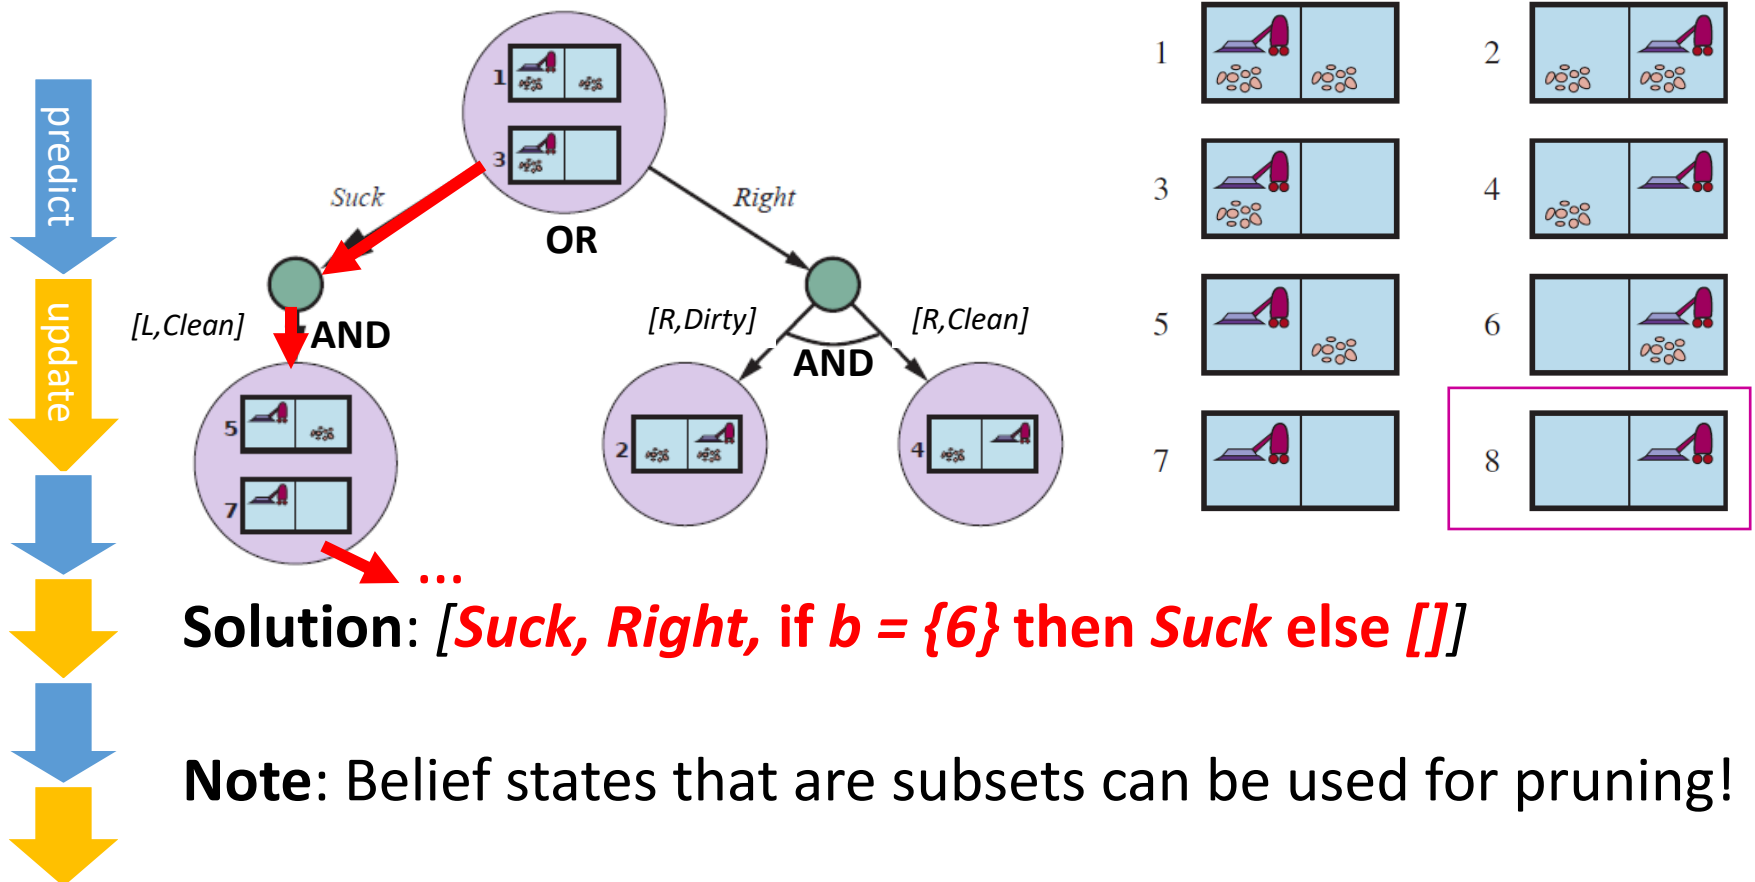
Solving Partially Observable Problems

Use an AND-OR tree to create a conditional plan



Solving Partially Observable Problems

Use an AND-OR tree to create a conditional plan



State Estimation and Approximate Belief States


- Agents choose an **action** and then receive an **observation** from the environment.
- The agent keep track of its belief state using the following update:

$$b \leftarrow \text{Update}(\text{Predict}(b, a), o)$$

- This process is often called
 - **monitoring**,
 - **filtering**, or
 - **state estimation**.
- The agent needs to be able to update its belief state following observations in **real time**! For many practical application, there is only time to compute an **approximate belief state**! These approximate methods are outside the scope of this introductory course.

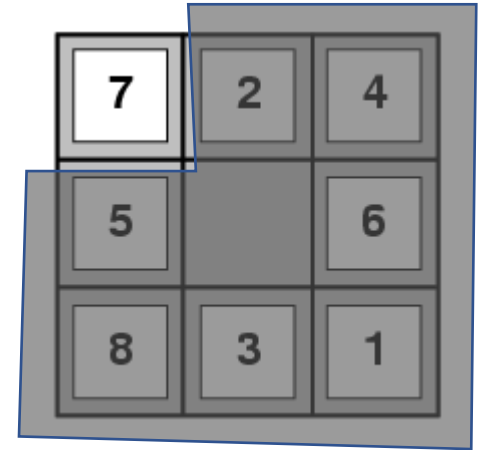
Case Study

Partially Observable 8-Puzzle



Partially Observable 8-Puzzle

- How do we solve this problem? What are the main steps?
- Give a problem description for each step.
 - States:
 - Initial state:
 - Actions:
 - Transition model:
 - Goal test:
 - Percept function:
- What algorithms can be used?





Unknown Environment

Online Search

Online Search

- **Recall offline search:** Create plan using the state space as a model before taking any action. The plan can be a sequence of actions or a conditional plan (for all possible observations).
- **Online search** explores the real world one action at a time. Prediction is replaced by “act” and update by “observe.”



- Useful for
 - **Real-time problems:** When offline computation takes too long and there is a penalty for sitting around and thinking.
 - **Nondeterministic domain:** Deal with what actually happens.
 - **Unknown environment:** The agent needs to explore to map an unknown area (state space).
The map is the transition model $f : S \times A \rightarrow S$

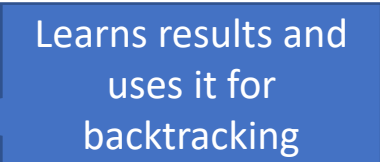
Issues With Online Search


- **Knowledge:** What does the agent know about the outcome of actions? E.g., does go north and then south lead to the same location?
- **Safely explorable state space/world:** There are **no irreversible actions** that cannot be undone (e.g., traps, cliffs). At least the agent does not execute these actions.
- Expanding nodes in **local order** is more efficient:
Depth-first search with back tracking

DFS Online Search Algorithm

Environment is deterministic and completely observable ($\text{percept}(s) = s'$ after action a), but the transition model (function result) is unknown. This algorithm builds the map result by trying all actions and backtracks when all actions in a state have been explored.

```
function ONLINE-DFS-AGENT(problem,  $s'$ ) returns an action  
     $s$ ,  $a$ , the previous state and action, initially null  
    persistent:  $\text{result}$ , a table mapping  $(s, a)$  to  $s'$ , initially empty  
                 $\text{untried}$ , a table mapping  $s$  to a list of untried actions  
                 $\text{unbacktracked}$ , a table mapping  $s$  to a list of states never backtracked to  
  
    if problem.IS-GOAL( $s'$ ) then return stop  
    if  $s'$  is a new state (not in  $\text{untried}$ ) then  $\text{untried}[s'] \leftarrow \text{problem}.\text{ACTIONS}(s')$   
    if  $s$  is not null then  
         $\text{result}[s, a] \leftarrow s'$   
        add  $s$  to the front of  $\text{unbacktracked}[s']$  // unbacktraced leaves a “breadcrumb trail”  
    if  $\text{untried}[s']$  is empty then  
        if  $\text{unbacktracked}[s']$  is empty then return stop  
        else  $a \leftarrow$  an action  $b$  such that  $\text{result}[s', b] = \text{POP}(\text{unbacktracked}[s'])$   
    else  $a \leftarrow \text{POP}(\text{untried}[s'])$   
     $s \leftarrow s'$   
    return  $a$ 
```





Important concepts that you should be able to explain and use now...

- Difference between the solution types:
 - a fixed actions sequence, and
 - a conditional plan (also called strategy or policy).
- What are belief states?
- How actions can be used to coerce the world into states.
- State estimation with repeated predict and update steps.
- The use of AND-OR trees.